



Structured extraction models: consistency, speed, size, and cost

30 NHTSA recall descriptions × 5 repeats per model, temperature 0; each panel ranked best to worst


Consistency: % of documents with all three fields identical across 5 repeats (higher is better)




Speed: mean seconds per extraction [min—max] (quicker is better; dark inner bar = mean excl. failures)



Size on disk: Ollama download size (smaller is better)



API price: USD per 1M tokens, input + output (\$0 is preferred)



Note: Consistency is a very conservative calculation that primarily leverages the model's consistent interpretation of the extraction scheme: identical repeats do not mean the model abided by the scheme, and this is not necessarily a measure of accuracy. This is purely the Gage's Repeatability. **Hardware:** Runtimes were measured on a Windows 11 laptop with an 11th Gen Intel Core i7-1185G7 @ 3.00 GHz, 16 GB RAM, and a standard 128 MB Intel Iris Xe graphics card (CPU-only inference). **Software:** R 4.6.0 with ellmer 0.4.1; gpt-5.4-nano via the OpenAI API. Run on July 12, 2026.