



Structured extraction models: consistency, speed, size, and cost

30 NHTSA recall descriptions × 5 repeats per model, temperature 0; each panel ranked best to worst

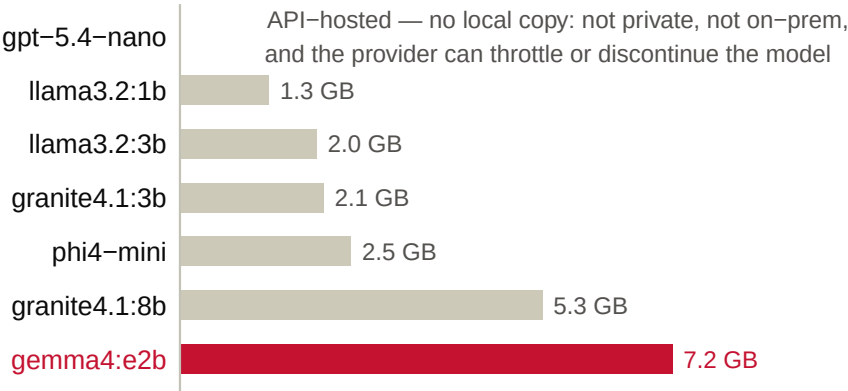
Consistency: % of documents with all three fields identical across 5 repeats (higher is better)



Speed: mean seconds per extraction [min—max] (quicker is better; dark inner bar = mean excl. failures)



Size on disk: Ollama download size (smaller is better)



API price: USD per 1M tokens, input + output (\$0 is preferred)



Note: Consistency is a very conservative calculation that primarily leverages the model's consistent interpretation of the extraction scheme: identical repeats do not mean the model abided by the scheme, and this is not necessarily a measure of accuracy. This is purely the Gage's Repeatability. **Hardware:** Runtimes were measured on a Windows 11 desktop with an Intel Core i9-14900K @ 3.20 GHz, 64 GB RAM, and an NVIDIA RTX 5000 Ada GPU with 32 GB VRAM (GPU-accelerated local inference). **Software:** R 4.6.1 with ellmer 0.4.1.9000; local models served via Ollama 0.31.2; gpt-5.4-nano via the OpenAI API. Run on July 12–13, 2026.